

Final Project 2 - Diabetic Retinopathy Classification

April 28, 2025

1 1. Importing Libraries

```
[98]: #Basic libraries
import os
import numpy as np
import tensorflow as tf
import pandas as pd

#Visualization
import matplotlib.pyplot as plt
import seaborn as sns

#Tensor Flow Packages
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras import layers, models

#DL/CNN Model
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.applications import VGG16
from tensorflow.keras.applications import InceptionV3

#ML Classification Model
from sklearn.ensemble import RandomForestClassifier

#Metrics
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score, classification_report
```

2 2. Loading and Preprocessing Images

```
[100]: # Set up directories for images
image_dir = 'Desktop/My Courses/Spring 2025/Machine Learning/Final Project/
↳retinopathy_3Classes' # path to images
labels = [] # List to hold labels (0, 1, 2 for healthy, moderate DR, severe DR)
image_paths = [] # List to hold image paths
```

```

# Collect image paths and corresponding labels
for filename in os.listdir(image_dir):
    if filename.endswith('.jpeg'):
        id_side_class = filename.split('_') # ID_Side_Class.jpeg
        label = int(id_side_class[-1].split('.')[0]) # Last part is the class
        image_paths.append(os.path.join(image_dir, filename))
        labels.append(label)

# Convert labels to numpy array
labels = np.array(labels)

# Load and preprocess images
images = []
for path in image_paths:
    img = tf.keras.preprocessing.image.load_img(path, target_size=(224, 224))
    img_array = tf.keras.preprocessing.image.img_to_array(img) / 255.0 #
    ↪ Normalize
    images.append(img_array)

images = np.array(images)

# Split into training and testing sets (80% / 20%)
X_train, X_test, y_train, y_test = train_test_split(images, labels, test_size=0.
    ↪ 2, stratify=labels, random_state=42)

```

3 3. Build Models

3.1 i. ResNet-50 Model

```
[103]: # Load ResNet50 model
base_model_resnet = ResNet50(weights='imagenet', include_top=False,
    ↪input_shape=(224, 224, 3))

# Freeze the base model layers
base_model_resnet.trainable = False

# Create the model
model_resnet = models.Sequential([
    base_model_resnet, # Pre-trained ResNet50 base
    layers.GlobalAveragePooling2D(), # Global average pooling to reduce the
    ↪output to a vector
    layers.Dense(3, activation='softmax') # Output layer with 3 classes
])

# Compile the model
model_resnet.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
    ↪metrics=['accuracy'])
```

3.2 ii. VGG-16 Model

```
[105]: # Load VGG-16 model
base_model_vgg = VGG16(weights='imagenet', include_top=False, input_shape=(224,
    ↪224, 3))

# Freeze the base model layers
base_model_vgg.trainable = False

# Create the model
model_vgg = models.Sequential([
    base_model_vgg, # Pre-trained VGG16 base
    layers.GlobalAveragePooling2D(), # Global average pooling to reduce the
    ↪output to a vector
    layers.Dense(3, activation='softmax') # Output layer with 3 classes
])

# Compile the model
model_vgg.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
    ↪metrics=['accuracy'])
```

3.3 iii. InceptionV3 Model

```
[107]: # Load InceptionV3 model
base_model_inception = InceptionV3(weights='imagenet', include_top=False,
    ↪input_shape=(224, 224, 3))

# Freeze the base model layers
```

```

base_model_inception.trainable = False

# Create the model
model_inception = models.Sequential([
    base_model_inception, # Pre-trained InceptionV3 base
    layers.GlobalAveragePooling2D(), # Global average pooling to reduce the
    ↪output to a vector
    layers.Dense(3, activation='softmax') # Output layer with 3 classes
])

# Compile the model
model_inception.compile(optimizer='adam',
    ↪loss='sparse_categorical_crossentropy', metrics=['accuracy'])

```

4 4. ML Classification Model

```

[109]: # Flatten the images so they can be used in traditional ML models (Random
    ↪Forest)
X_train_flattened = X_train.reshape(X_train.shape[0], -1) # Flatten the images
    ↪into 1D vectors
X_test_flattened = X_test.reshape(X_test.shape[0], -1)

# Train a Random Forest model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train_flattened, y_train)

# Predictions on the test set
rf_pred = rf_model.predict(X_test_flattened)

# Calculate Accuracy and F1-score
rf_accuracy = accuracy_score(y_test, rf_pred)
rf_f1 = f1_score(y_test, rf_pred, average='weighted')

print(f"Random Forest Accuracy: {rf_accuracy:.4f}")
print(f"Random Forest F1 Score: {rf_f1:.4f}")

```

Random Forest Accuracy: 0.5837

Random Forest F1 Score: 0.4876

5 5. Training Models

```

[111]: # Train the ResNet-50 model
history_resnet = model_resnet.fit(X_train, y_train, validation_data=(X_test,
    ↪y_test), epochs=10, batch_size=32)

# Train the VGG-16 model

```

```
history_vgg = model_vgg.fit(X_train, y_train, validation_data=(X_test, y_test),  
    ↪ epochs=10, batch_size=32)
```

```
# Train the InceptionV3 model
```

```
history_inception = model_inception.fit(X_train, y_train,  
    ↪ validation_data=(X_test, y_test), epochs=10, batch_size=32)
```

Epoch 1/10

53/53 50s 900ms/step -

accuracy: 0.5361 - loss: 0.9871 - val_accuracy: 0.5837 - val_loss: 0.9479

Epoch 2/10

53/53 44s 836ms/step -

accuracy: 0.5895 - loss: 0.9496 - val_accuracy: 0.5837 - val_loss: 0.9529

Epoch 3/10

53/53 44s 839ms/step -

accuracy: 0.5972 - loss: 0.9525 - val_accuracy: 0.5837 - val_loss: 0.9503

Epoch 4/10

53/53 44s 832ms/step -

accuracy: 0.5843 - loss: 0.9612 - val_accuracy: 0.5837 - val_loss: 0.9480

Epoch 5/10

53/53 55s 1s/step -

accuracy: 0.5806 - loss: 0.9669 - val_accuracy: 0.5837 - val_loss: 0.9457

Epoch 6/10

53/53 44s 840ms/step -

accuracy: 0.5808 - loss: 0.9619 - val_accuracy: 0.5837 - val_loss: 0.9513

Epoch 7/10

53/53 45s 860ms/step -

accuracy: 0.5896 - loss: 0.9510 - val_accuracy: 0.5837 - val_loss: 0.9582

Epoch 8/10

53/53 45s 853ms/step -

accuracy: 0.5601 - loss: 0.9616 - val_accuracy: 0.5837 - val_loss: 0.9459

Epoch 9/10

53/53 49s 921ms/step -

accuracy: 0.6046 - loss: 0.9281 - val_accuracy: 0.5837 - val_loss: 0.9472

Epoch 10/10

53/53 47s 888ms/step -

accuracy: 0.5822 - loss: 0.9601 - val_accuracy: 0.5837 - val_loss: 0.9496

Epoch 1/10

53/53 161s 3s/step -

accuracy: 0.5941 - loss: 0.9398 - val_accuracy: 0.5837 - val_loss: 0.9419

Epoch 2/10

53/53 160s 3s/step -

accuracy: 0.5990 - loss: 0.9220 - val_accuracy: 0.5837 - val_loss: 0.9214

Epoch 3/10

53/53 161s 3s/step -

accuracy: 0.5749 - loss: 0.9310 - val_accuracy: 0.5766 - val_loss: 0.9261

Epoch 4/10

53/53 154s 3s/step -

accuracy: 0.5854 - loss: 0.9111 - val_accuracy: 0.5837 - val_loss: 0.9077
 Epoch 5/10
 53/53 157s 3s/step -
 accuracy: 0.5852 - loss: 0.9153 - val_accuracy: 0.5813 - val_loss: 0.9144
 Epoch 6/10
 53/53 153s 3s/step -
 accuracy: 0.5730 - loss: 0.9146 - val_accuracy: 0.5837 - val_loss: 0.8992
 Epoch 7/10
 53/53 154s 3s/step -
 accuracy: 0.5847 - loss: 0.8910 - val_accuracy: 0.5837 - val_loss: 0.8967
 Epoch 8/10
 53/53 163s 3s/step -
 accuracy: 0.5856 - loss: 0.8977 - val_accuracy: 0.5837 - val_loss: 0.8938
 Epoch 9/10
 53/53 156s 3s/step -
 accuracy: 0.5856 - loss: 0.8892 - val_accuracy: 0.5837 - val_loss: 0.8916
 Epoch 10/10
 53/53 163s 3s/step -
 accuracy: 0.5854 - loss: 0.8987 - val_accuracy: 0.5861 - val_loss: 0.8906
 Epoch 1/10
 53/53 37s 655ms/step -
 accuracy: 0.5226 - loss: 1.0444 - val_accuracy: 0.5933 - val_loss: 0.8864
 Epoch 2/10
 53/53 34s 639ms/step -
 accuracy: 0.6023 - loss: 0.8548 - val_accuracy: 0.6268 - val_loss: 0.9165
 Epoch 3/10
 53/53 36s 675ms/step -
 accuracy: 0.6290 - loss: 0.8201 - val_accuracy: 0.6268 - val_loss: 0.8429
 Epoch 4/10
 53/53 34s 634ms/step -
 accuracy: 0.6831 - loss: 0.7238 - val_accuracy: 0.6364 - val_loss: 0.8334
 Epoch 5/10
 53/53 33s 627ms/step -
 accuracy: 0.6982 - loss: 0.7001 - val_accuracy: 0.5861 - val_loss: 0.8628
 Epoch 6/10
 53/53 33s 619ms/step -
 accuracy: 0.7064 - loss: 0.6859 - val_accuracy: 0.5789 - val_loss: 0.8776
 Epoch 7/10
 53/53 33s 629ms/step -
 accuracy: 0.7171 - loss: 0.6681 - val_accuracy: 0.6077 - val_loss: 0.8487
 Epoch 8/10
 53/53 36s 677ms/step -
 accuracy: 0.6907 - loss: 0.6990 - val_accuracy: 0.5933 - val_loss: 0.8604
 Epoch 9/10
 53/53 34s 639ms/step -
 accuracy: 0.7062 - loss: 0.6734 - val_accuracy: 0.6483 - val_loss: 0.8220
 Epoch 10/10
 53/53 32s 610ms/step -

accuracy: 0.7302 - loss: 0.6337 - val_accuracy: 0.6459 - val_loss: 0.8146

6. Evaluation and Metrics

```
[113]: # Get Accuracy and F1-Score for each model
resnet_acc = history_resnet.history['val_accuracy'][-1]
vgg_acc = history_vgg.history['val_accuracy'][-1]
inception_acc = history_inception.history['val_accuracy'][-1]

# Random Forest metrics computed above
rf_acc = rf_accuracy
rf_f1 = rf_f1

# Create a summary table with F1-scores computed
results = {
    'Model': ['ResNet-50', 'VGG-16', 'InceptionV3', 'Random Forest'],
    'Accuracy': [resnet_acc, vgg_acc, inception_acc, rf_acc],
    'F1 Score': [
        f1_score(y_test, model_resnet.predict(X_test).argmax(axis=1),
        ↪average='weighted'),
        f1_score(y_test, model_vgg.predict(X_test).argmax(axis=1),
        ↪average='weighted'),
        f1_score(y_test, model_inception.predict(X_test).argmax(axis=1),
        ↪average='weighted'),
        rf_f1
    ]
}

# Convert to DataFrame for easy viewing
results_df = pd.DataFrame(results)
print(results_df)
```

```
14/14          11s 730ms/step
14/14          32s 2s/step
14/14          10s 639ms/step

   Model  Accuracy  F1 Score
0  ResNet-50  0.583732  0.430304
1   VGG-16  0.586124  0.436243
2  InceptionV3  0.645933  0.613259
3  Random Forest  0.583732  0.487595
```

7. Visualizing Scores

```
[115]: # Reshape the DataFrame for easier plotting
results_df_melted = results_df.melt(id_vars="Model", value_vars=["Accuracy",
    ↪"F1 Score"],
                                     var_name="Metric", value_name="Score")
```


```

# Create the combined bar plot
plt.figure(figsize=(8, 6))
sns.barplot(x="Model", y="Score", hue="Metric", data=results_df_melted,
            palette="Set2")

# Titles and labels
plt.title('Comparison of Accuracy and F1-Score for Each Model')
plt.xticks(rotation=45)
plt.ylim(0, 1)
plt.ylabel('Score')

# Show plot
plt.tight_layout()
plt.show()

```



8. Classification Reports for each Model

```
[129]: # Classification Reports for each model
print("Classification Report for ResNet-50:") #ResNet-50
print(classification_report(y_test, model_resnet.predict(X_test).
    ↪argmax(axis=1), zero_division=0))

print("Classification Report for VGG-16:") #VGG-16
print(classification_report(y_test, model_vgg.predict(X_test).argmax(axis=1),
    ↪zero_division=0))

print("Classification Report for InceptionV3:") #InceptionV3
print(classification_report(y_test, model_inception.predict(X_test).
    ↪argmax(axis=1), zero_division=0))

print("Classification Report for Random Forest:") #Random Forest
print(classification_report(y_test, rf_pred, zero_division=0))
```

Classification Report for ResNet-50:

```
14/14          9s 630ms/step
      precision    recall  f1-score   support

     0         0.58        1.00        0.74        244
     1         0.00        0.00        0.00        112
     2         0.00        0.00        0.00         62

 accuracy                   0.58        418
 macro avg         0.19        0.33        0.25        418
 weighted avg      0.34        0.58        0.43        418
```

Classification Report for VGG-16:

```
14/14          33s 2s/step
      precision    recall  f1-score   support

     0         0.59        1.00        0.74        244
     1         0.00        0.00        0.00        112
     2         0.50        0.02        0.03         62

 accuracy                   0.59        418
 macro avg         0.36        0.34        0.26        418
 weighted avg      0.42        0.59        0.44        418
```

Classification Report for InceptionV3:

```
14/14          7s 466ms/step
      precision    recall  f1-score   support

     0         0.70        0.86        0.77        244
```

1	0.44	0.21	0.29	112
2	0.57	0.60	0.58	62
accuracy			0.65	418
macro avg	0.57	0.56	0.55	418
weighted avg	0.61	0.65	0.61	418

Classification Report for Random Forest:

	precision	recall	f1-score	support
0	0.61	0.94	0.74	244
1	0.33	0.12	0.17	112
2	0.40	0.03	0.06	62
accuracy			0.58	418
macro avg	0.45	0.36	0.32	418
weighted avg	0.51	0.58	0.49	418

[]: